

Fluid & JuiceFS & Dragonfly-EN V2

Architecture

Install JuiceFS with Dragonfly

Dragonfly Kubernetes Cluster Setup

Setup Kubernetes Cluster

Kind loads dragonfly image

Create dragonfly cluster based on helm charts

Expose Dragonfly Peer Proxy service port

Install JuiceFS

Verify

Multi-Node Read Performance Testing

JuiceFS

JuiceFS & Dragonfly

Analysis

Install Fluid & JuiceFS Runtime with Dragonfly

Dragonfly Kubernetes Cluster Setup

Setup Kubernetes Cluster

Kind loads dragonfly image

Create dragonfly cluster based on helm charts

Expose Dragonfly Peer Proxy service port

Install Fluid

Create Fluid cluster based on helm charts

Create Dataset

Create JuiceFS Runtime

Verify

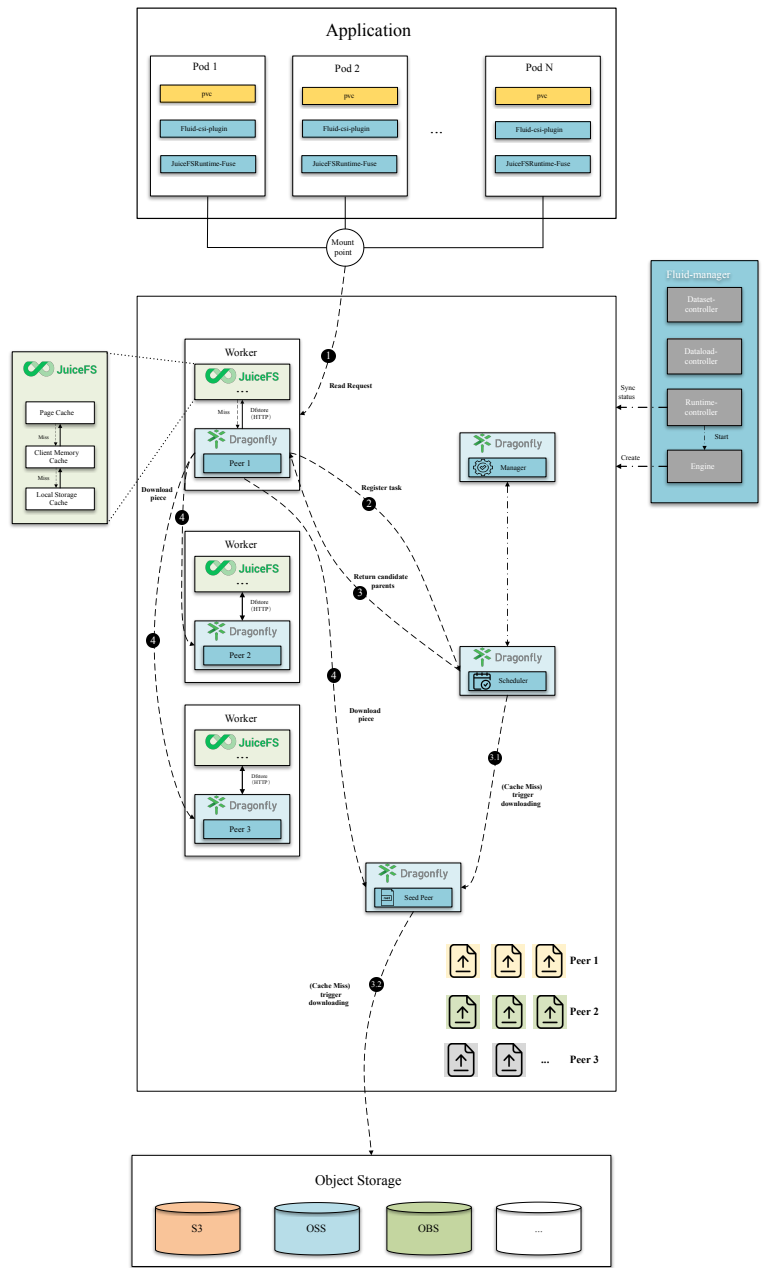
Reference

This document will help you experience how to use dragonfly with [Fluid](#) & [JuiceFS Runtime](#). It introduces how to install Fluid & JuiceFS Runtime with Dragonfly and performance

testing.

Fluid supports Alluxio Runtime, JuiceFS Runtime, Thin Runtime, etc. However, Fluid still has scope for optimization for large file downloads at the same time. For example, Facebook LLaMa 65B which is 121.62 GB. If use the memory cache or disk cache on a single node, It is a hard work for the storage and network in the machine. [Dragonfly](#) Provides efficient, stable, secure file distribution and image acceleration based on p2p technology. It is designed to improve the efficiency and speed of large-scale file distribution. **JuiceFS depend on Dragonfly for object storage, and realizes the integration of Fluid and Dragonfly through JuiceFS Runtime..**

Architecture



Dragonfly becomes a new cache between JuiceFS and object storage. There are optimizations in the reading and writing. When reading, if there is no hit in the JuiceFS cache, the traffic will be forwarded to Dragonfly Peer. It can be used to eliminate the bandwidth limit of the object storage through P2P technology, thereby accelerating file downloading.

Install JuiceFS with Dragonfly

Dragonfly Kubernetes Cluster Setup

Setup Kubernetes Cluster

[Kind](#) is recommended if no Kubernetes cluster is available for testing.

Create kind multi-node cluster configuration file `kind-config.yaml`, configuration content is as follows:

```
▼ Shell |
1  kind: Cluster
2  apiVersion: kind.x-k8s.io/v1alpha4
3  nodes:
4    - role: control-plane
5    - role: worker
6    - role: worker
```

Create a kind multi-node cluster using the configuration file:

```
▼ Shell |
1  kind create cluster --config kind-config.yaml
```

Kind loads dragonfly image

Pull dragonfly latest images:

```
▼ Shell |
1  docker pull dragonflyoss/scheduler:latest
2  docker pull dragonflyoss/manager:latest
3  docker pull dragonflyoss/dfdaemon:latest
```

Kind cluster loads dragonfly latest images:

```
▼ Shell |
1  kind load docker-image dragonflyoss/scheduler:latest
2  kind load docker-image dragonflyoss/manager:latest
3  kind load docker-image dragonflyoss/dfdaemon:latest
```

Create dragonfly cluster based on helm charts

Create helm charts configuration file `charts-config.yaml` and set `dfdaemon.config.proxy.proxies` and `seedPeer.config.proxy.proxies` to math endpoint of the object storage, configuration content is as follows:

```
1 scheduler:
2   image: dragonflyoss/scheduler
3   tag: latest
4   replicas: 1
5   metrics:
6     enable: true
7   config:
8     verbose: true
9     pprofPort: 18066
10
11 seedPeer:
12   image: dragonflyoss/dfdaemon
13   tag: latest
14   replicas: 1
15   metrics:
16     enable: true
17   config:
18     verbose: true
19     pprofPort: 18066
20   proxy:
21     defaultFilter: 'Expires&Signature&ns'
22     security:
23       insecure: true
24       tlsVerify: false
25     tcpListen:
26       # # Listen address.
27       # listen: 0.0.0.0
28       # Listen port, daemon will try to listen,
29       # when this port is not available, daemon will try next port.
30       port: 65001
31       namespace: ""
32     proxies:
33       # Proxy all http download requests of the s3.
34       - regx: s3.*amazonaws.com.*
35       # Proxy all http download requests of the oss.
36       - regx: oss.*aliyuncs.com.*
37       # Proxy all http download requests of the obs.
38       - regx: obs.*myhuaweicloud.com.*
39
40 dfdaemon:
41   image: dragonflyoss/dfdaemon
42   tag: latest
43   metrics:
44     enable: true
45   config:
```

```

46     verbose: true
47     pprofPort: 18066
48     proxy:
49         defaultFilter: 'Expires&Signature&ns'
50         security:
51             insecure: true
52             tlsVerify: false
53         tcpListen:
54             # # Listen address.
55             # listen: 0.0.0.0
56             # Listen port, daemon will try to listen,
57             # when this port is not available, daemon will try next port.
58             port: 65001
59             namespace: ""
60         proxies:
61             # Proxy all http download requests of the s3.
62             - regx: s3.*amazonaws.com.*
63             # Proxy all http download requests of the oss.
64             - regx: oss.*aliyuncs.com.*
65             # Proxy all http download requests of the obs.
66             - regx: obs.*myhuaweicloud.com.*
67
68     manager:
69         image: dragonflyoss/manager
70         tag: latest
71         replicas: 1
72         metrics:
73             enable: true
74         config:
75             verbose: true
76             pprofPort: 18066
77
78     jaeger:
79         enable: true

```

Create a dragonfly cluster using the configuration file:

```
1 $ helm repo add dragonfly https://dragonflyoss.github.io/helm-charts/
2 $ helm install --wait --create-namespace --namespace dragonfly-system drag
  onfly dragonfly/dragonfly -f charts-config.yaml
3 NAME: dragonfly
4 LAST DEPLOYED: Thu Sep 28 17:35:49 2023
5 NAMESPACE: dragonfly-system
6 STATUS: deployed
7 REVISION: 1
8 TEST SUITE: None
9 NOTES:
10 1. Get the scheduler address by running these commands:
11   export SCHEDULER_POD_NAME=$(kubectl get pods --namespace dragonfly-syste
  m -l "app=dragonfly,release=dragonfly,component=scheduler" -o jsonpath={.i
  tems[0].metadata.name})
12   export SCHEDULER_CONTAINER_PORT=$(kubectl get pod --namespace dragonfly-
  system $SCHEDULER_POD_NAME -o jsonpath="{.spec.containers[0].ports[0].cont
  ainerPort}")
13   kubectl --namespace dragonfly-system port-forward $SCHEDULER_POD_NAME 80
  02:$SCHEDULER_CONTAINER_PORT
14   echo "Visit http://127.0.0.1:8002 to use your scheduler"
15
16 2. Get the dfdaemon port by running these commands:
17   export DFDAEMON_POD_NAME=$(kubectl get pods --namespace dragonfly-syste
  m -l "app=dragonfly,release=dragonfly,component=dfdaemon" -o jsonpath={.it
  ems[0].metadata.name})
18   export DFDAEMON_CONTAINER_PORT=$(kubectl get pod --namespace dragonfly-s
  ystem $DFDAEMON_POD_NAME -o jsonpath="{.spec.containers[0].ports[0].contai
  nerPort}")
19   You can use $DFDAEMON_CONTAINER_PORT as a proxy port in Node.
20
21 3. Configure runtime to use dragonfly:
22   https://d7y.io/docs/getting-started/quick-start/kubernetes/
23
24
25 4. Get Jaeger query URL by running these commands:
26   export JAEGER_QUERY_PORT=$(kubectl --namespace dragonfly-system get serv
  ices dragonfly-jaeger-query -o jsonpath="{.spec.ports[0].port}")
27   kubectl --namespace dragonfly-system port-forward service/dragonfly-jaeg
  er-query 16686:$JAEGER_QUERY_PORT
28   echo "Visit http://127.0.0.1:16686/search?limit=20&lookback=1h&maxDurati
  on&minDuration&service=dragonfly to query download events"
```

Check that dragonfly is deployed successfully:


```
Shell |
1 $ kubectl get po -n dragonfly-system
2 NAME                                READY   STATUS    RESTARTS   AGE
3 dragonfly-dfdaemon-65rz7            1/1    Running   5 (6m17s ago)  8m4
4   3s
5 dragonfly-dfdaemon-rnvsj            1/1    Running   5 (6m23s ago)  8m4
6   3s
7 dragonfly-jaeger-7d58dfcfc8-qmn8c  1/1    Running   0           8m4
8   3s
9 dragonfly-manager-6f8b4f5c66-qq8sd 1/1    Running   0           8m4
10  3s
11 dragonfly-mysql-0                  1/1    Running   0           8m4
12  3s
13 dragonfly-redis-master-0           1/1    Running   0           8m4
14  3s
15 dragonfly-redis-replicas-0         1/1    Running   0           8m4
16  3s
17 dragonfly-redis-replicas-1         1/1    Running   0           7m3
18  3s
19 dragonfly-redis-replicas-2         1/1    Running   0           5m5
20  0s
21 dragonfly-scheduler-0              1/1    Running   0           8m4
22  3s
23 dragonfly-seed-peer-0              1/1    Running   3 (5m56s ago) 8m4
24  3s
```

Expose Dragonfly Peer Proxy service port

Create the `proxy.yaml` configuration to expose the port on which the Dragonfly Peer Proxy listens. The default port is `65001` and set `targetPort` to `65001`.

```
YAML |
1 kind: Service
2 apiVersion: v1
3 metadata:
4   name: proxy
5 spec:
6   selector:
7     app: dragonfly
8     component: dfdaemon
9     release: dragonfly
10
11   ports:
12   - protocol: TCP
13     port: 65001
14     targetPort: 65001
15
16   type: NodePort
```

Create service:

```
Plain Text |
1 kubectl --namespace dragonfly-system apply -f proxy.yaml
```

Forward request to Dragonfly Peer Proxy:

```
Shell |
1 kubectl --namespace dragonfly-system port-forward service/proxy 65001:65001
```

Install JuiceFS

For detailed installation documentation, please refer to [JuiceFS document](#). For Linux and macOS systems, you can use a one-click installation script that automatically downloads and installs the latest version of the JuiceFS client based on your hardware architecture.

```
Shell |
1 # 默认安装到 /usr/local/bin
2 curl -sSL https://d.juicefs.com/install | sh -
```

After installation, you can specify the use of Dragonfly as the object storage when executing commands such as `juicefs format` and `juicefs config` :

```
1 juicefs format \  
2   --storage dragonfly \  
3   --access-key ABCDEFGHIJKLMNOPqXYZ \  
4   --secret-key ZYXwvutsrqpoNMLkJiHgfeDCBA \  
5   --bucket "https://myjfs.oss-cn-shanghai.aliyuncs.com?proxy=http://127.  
6     0.0.1:65001&backendStorage=oss" \  
7     redis://192.168.1.6:6379/1 \  
8     myjfs
```

`--storage` : Set the object storage type, the value is `dragonfly` .

`--bucket` : Set the object storage endpoint.

`--access-key` : Set the object storage access key.

`--secret-key` : Set the object storage secret key.

Dragonfly Peer Proxy address and Backend object storage type are added to the query string in the `--bucket` . The details of parameter is as follows:

Param	Type	Describe	Required
backendStorage	string	Backend object storage type, supports s3, oss and obs.	Y
proxy	string	Draognfly Peer Proxy address.	Y

Verify the created file system status:

```

1  $ juicefs status redis://192.168.1.6:6379/1
2  2023/10/17 19:09:35.738635 juicefs[2273224] <INFO>: Meta address: redis://
   localhost:6379/1 [interface.go:498]
3  2023/10/17 19:09:35.739344 juicefs[2273224] <WARNING>: AOF is not enable
   d, you may lose data if Redis is not shutdown properly. [info.go:84]
4  2023/10/17 19:09:35.739407 juicefs[2273224] <INFO>: Ping redis latency: 22
   .384µs [redis.go:3572]
5  {
6  "Setting": {
7      "Name": "myjfs",
8      "UUID": "316d39df-a7ba-4cde-8cc7-5568a7a0f745",
9      "Storage": "dragonfly",
10     "Bucket": "https://myjfs.oss-cn-shanghai.aliyuncs.com?proxy=http://12
       7.0.0.1:65001\u0026backendStorage=oss",
11     "BlockSize": 4096,
12     "Compression": "none",
13     "EncryptAlgo": "aes256gcm-rsa",
14     "TrashDays": 1,
15     "MetaVersion": 1,
16     "MinClientVersion": "1.1.0-A",
17     "DirStats": true
18 },
19 "Sessions": [],
20 "Statistic": {
21     "UsedSpace": 0,
22     "AvailableSpace": 1125899906842624,
23     "UsedInodes": 0,
24     "AvailableInodes": 10485760
25 }
26 }
27

```

When using other JuiceFS commands, you can also specify Dragonfly as the object storage. For detailed JuiceFS commands documentation, please refer to [document](#).

Verify

Execute the command:

```
1 juicefs objbench \  
2   --storage dragonfly \  
3   --access-key ABCDEFGHIJKLMNOPqXYZ \  
4   --secret-key ZYXwvutsrqpoNMLkjiHgfeDCBA \  
5   https://myjfs.oss-cn-shanghai.aliyuncs.com?proxy=http://127.0.0.1:65001  
   &backendStorage=oss
```

The endpoint is `https://myjfs.oss-cn-shanghai.aliyuncs.com?proxy=http://127.0.0.1:65001&backendStorage=oss`. It should pass the unit test for `teststorage`.

```

1  +-----+-----+-----+-----+
2  | CATEGORY |          TEST          |          RESULT          |
3  +-----+-----+-----+-----+
4  |  basic  | create a bucket |
   | pass   |
5  |  basic  |   put an object |
   | pass   |
6  |  basic  |   get an object |
   | pass   |
7  |  basic  |   get non-exist |
   | pass   |
8  |  basic  | get partial object | failed to get object with the offset ou
   | t of r... |
9  |  basic  |   head an object |
   | pass   |
10 |  basic  | delete an object |
   | pass   |
11 |  basic  | delete non-exist |
   | pass   |
12 |  basic  |   list objects |
   | pass   |
13 |  basic  |   special key |
   | pass   |
14 |  sync  | put a big object |
   | pass   |
15 |  sync  | put an empty object |
   | pass   |
16 |  sync  | multipart upload |
   | pass   |
17 |  sync  | change owner/group |
   | t support |
18 |  sync  | change permission |
   | t support |
19 |  sync  |   change mtime |
   | t support |
20 +-----+-----+-----+-----+
    +-----+

```

Multi-Node Read Performance Testing

JuiceFS

Test the caching performance of JuiceFS. The configured object storage needs to be the same as in Dragonfly.

```
1 juicefs format \  
2   --storage oss \  
3   --bucket https://myjfs.oss-cn-shanghai.aliyuncs.com \  
4   --access-key ABCDEFGHIJKLMNopqXYZ \  
5   --secret-key ZYXwvutsrqpoNMLkJiHgfeDCBA \  
6   redis://192.168.1.6:6379/2 \  
7   myjfs
```

Mount the file system using the `juicefs mount` command:

```
1 juicefs mount redis://192.168.1.6:6379/2 /mnt/jfs
```

Creat a 1GB file in the mounted directory:

```
1 $ time dd if=/dev/zero of=/mnt/jfs/test.txt bs=1M count=1000  
2 1000+0 records in  
3 1000+0 records out  
4 1048576000 bytes (1.0 GB, 1000 MiB) copied, 10.7013 s, 98.0 MB/s  
5 dd if=/dev/zero of=/mnt/jfs/test.txt bs=1M count=1000 0.00s user 0.33s sys  
   tem 3% cpu 10.711 tota
```

For the first read, JuiceFS triggers **back-to-source download** and it takes **11.356 seconds**.

```
1 $ time cp /mnt/jfs/test.txt /dev/null  
2 cp /mnt/jfs/test.txt /dev/null 0.00s user 0.29s system 2% cpu 11.356 total
```

Clear the page cache and read again. Hit JuiceFS's cache, and it takes **0.347 seconds**.

```

1 $ sync && echo 3 > /proc/sys/vm/drop_caches
2 $ time cp /mnt/jfs/test.txt /dev/null
3 cp /mnt/jfs/test.txt /dev/null 0.00s user 0.30s system 86% cpu 0.347 total

```

JuiceFS & Dragonfly

Test the performance of Dragonfly cache and hit local peer cache and remote peer cache. Expose Draognfly Peer's `65001` port.

```

1 export dragonfly_dfdaemon_name=$(kubectl get po -n dragonfly-system | grep
dragonfly-dfdaemon- | tail -n 1 | awk '{print $1}')
2 kubectl --namespace dragonfly-system port-forward $dragonfly_dfdaemon_name
65001:65001

```

Initialize the file system based on Dragonfly:

```

1 juicefs format \
2   --storage dragonfly \
3   --access-key ABCDEFGHIJKLMNopqXYZ \
4   --secret-key ZYXwvutsrqpoNMLkJiHgfeDCBA \
5   --bucket "https://myjfs.oss-cn-shanghai.aliyuncs.com?proxy=http://127.
0.0.1:65001&backendStorage=oss" \
6   redis://192.168.1.6:6379/1 \
7   myjfs

```

Mount the file system and disable JuiceFS's cache:

```

1 juicefs mount redis://192.168.1.6:6379/1 /mnt/jfs --cache-size=0

```

Create a 1GB file in the mounted directory:


```
Shell |
1 $ time dd if=/dev/zero of=/mnt/jfs/test.txt bs=1M count=1000
2 1000+0 records in
3 1000+0 records out
4 1048576000 bytes (1.0 GB, 1000 MiB) copied, 10.2689 s, 102 MB/s
5 dd if=/dev/zero of=/mnt/jfs/test.txt bs=1M count=1000 0.00s user 0.38s sys
   tem 3% cpu 10.271 total
```

For the first read. No cache hits for JuiceFS and Dragonfly, and it triggers **back-to-source download**, taking **11.147 seconds**.

```
Shell |
1 $ time cp /mnt/jfs/test.txt /dev/null
2 cp /mnt/jfs/test.txt /dev/null 0.00s user 0.30s system 2% cpu 11.147 total
```

Clear the cache of the file system and read again. Hit the cache of **Dragonfly's Local Peer** and it takes **1.554 seconds**.

```
Shell |
1 $ sync && echo 3 > /proc/sys/vm/drop_caches
2 $ time cp /mnt/jfs/test.txt /dev/null
3 cp /mnt/jfs/test.txt /dev/null 0.00s user 0.32s system 20% cpu 1.554 total
```

Test the cache speed of the hit Dragonfly Remote Peer, delete the Peer:

```

Shell |
1 ▾ $ export dragonfly_dfdaemon_name=$(kubectl get po -n dragonfly-system --so
    rt-by=.metadata.creationTimestamp | grep "dragonfly-dfdaemon-" | awk '{pri
    nt $1}' | tail -n 1)
2 $ kubectl delete po $dragonfly_dfdaemon_name -n dragonfly-system
3 $ kubectl get po -n dragonfly-system
4 NAME                                READY    STATUS    RESTARTS    AGE
5 dragonfly-dfdaemon-5q4r8            1/1     Running  0           30s # ne
    w pod
6 dragonfly-dfdaemon-nhzcc            1/1     Running  0           19m
7 dragonfly-jaeger-c7947b579-q4hr4    1/1     Running  0           19m
8 dragonfly-manager-5dc5fbf548-zrf7d  1/1     Running  0           19m
9 dragonfly-mysql-0                  1/1     Running  0           19m
10 dragonfly-redis-master-0           1/1     Running  0           19m
11 dragonfly-redis-replicas-0         1/1     Running  0           19m
12 dragonfly-redis-replicas-1         1/1     Running  0           18m
13 dragonfly-redis-replicas-2         1/1     Running  0           18m
14 dragonfly-scheduler-0              1/1     Running  0           19m
15 dragonfly-seed-peer-0              1/1     Running  0           19m

```

Recreate the pod:

```

Shell |
1 kubectl --namespace dragonfly-system port-forward $dragonfly_dfdaemon_name
    65001:65001

```

Clear the cache of the file system and read again. The created Pod has no cache, and it hits the cache of the Remote Peer, it takes **1.937 seconds**.


```

Shell |
1 $ sync && echo 3 > /proc/sys/vm/drop_caches
2 $ time cp /mnt/jfs/test.txt /dev/null
3 cp /mnt/jfs/test.txt /dev/null 0.01s user 0.32s system 16% cpu 1.937 total

```

Analysis

Performance Testing



Test results show JuiceFS and Dragonfly integration. It can effectively reduce the file download time. Due to the influence of the network environment of the machine itself, the actual download time is not important, but the ratio of the increase in the download time in different scenarios is very important.

Install Fluid & JuiceFS Runtime with Dragonfly

Dragonfly Kubernetes Cluster Setup

Setup Kubernetes Cluster

[Kind](#) is recommended if no Kubernetes cluster is available for testing.

Create kind multi-node cluster configuration file `kind-config.yaml`, configuration content is as follows:

```
▼ Shell |
1 kind: Cluster
2 apiVersion: kind.x-k8s.io/v1alpha4
3 nodes:
4   - role: control-plane
5   - role: worker
6   - role: worker
```

Create a kind multi-node cluster using the configuration file:

```
▼ Shell |
1 kind create cluster --config kind-config.yaml
```

Kind loads dragonfly image

Pull dragonfly latest images:

```
▼ Shell |
1 docker pull dragonflyoss/scheduler:latest
2 docker pull dragonflyoss/manager:latest
3 docker pull dragonflyoss/dfdaemon:latest
```

Kind cluster loads dragonfly latest images:

```
▼ Shell |
1 kind load docker-image dragonflyoss/scheduler:latest
2 kind load docker-image dragonflyoss/manager:latest
3 kind load docker-image dragonflyoss/dfdaemon:latest
```

Create dragonfly cluster based on helm charts

Create helm charts configuration file `charts-config.yaml` and set `dfdaemon.config.proxy.proxies` and `seedPeer.config.proxy.proxies` to math endpoint of the object storage, configuration content is as follows:

```
1 scheduler:
2   image: dragonflyoss/scheduler
3   tag: latest
4   replicas: 1
5   metrics:
6     enable: true
7   config:
8     verbose: true
9     pprofPort: 18066
10
11 seedPeer:
12   image: dragonflyoss/dfdaemon
13   tag: latest
14   replicas: 1
15   metrics:
16     enable: true
17   config:
18     verbose: true
19     pprofPort: 18066
20   proxy:
21     defaultFilter: 'Expires&Signature&ns'
22     security:
23       insecure: true
24       tlsVerify: false
25     tcpListen:
26       # # Listen address.
27       # listen: 0.0.0.0
28       # Listen port, daemon will try to listen,
29       # when this port is not available, daemon will try next port.
30       port: 65001
31       namespace: ""
32     proxies:
33       # Proxy all http download requests of the s3.
34       - regx: s3.*amazonaws.com.*
35       # Proxy all http download requests of the oss.
36       - regx: oss.*aliyuncs.com.*
37       # Proxy all http download requests of the obs.
38       - regx: obs.*myhuaweicloud.com.*
39
40 dfdaemon:
41   image: dragonflyoss/dfdaemon
42   tag: latest
43   metrics:
44     enable: true
45   config:
```

```
46     verbose: true
47     pprofPort: 18066
48     proxy:
49       defaultFilter: 'Expires&Signature&ns'
50       security:
51         insecure: true
52         tlsVerify: false
53       tcpListen:
54         # # Listen address.
55         # listen: 0.0.0.0
56         # Listen port, daemon will try to listen,
57         # when this port is not available, daemon will try next port.
58         port: 65001
59         namespace: ""
60       proxies:
61         # Proxy all http download requests of the s3.
62         - regx: s3.*amazonaws.com.*
63         # Proxy all http download requests of the oss.
64         - regx: oss.*aliyuncs.com.*
65         # Proxy all http download requests of the obs.
66         - regx: obs.*myhuaweicloud.com.*
67
68     manager:
69       image: dragonflyoss/manager
70       tag: latest
71       replicas: 1
72     metrics:
73       enable: true
74     config:
75       verbose: true
76       pprofPort: 18066
77
78     jaeger:
79       enable: true
```

Create a dragonfly cluster using the configuration file:

```
1 $ helm repo add dragonfly https://dragonflyoss.github.io/helm-charts/
2 $ helm install --wait --create-namespace --namespace dragonfly-system drag
  onfly dragonfly/dragonfly -f charts-config.yaml
3 NAME: dragonfly
4 LAST DEPLOYED: Thu Sep 28 17:35:49 2023
5 NAMESPACE: dragonfly-system
6 STATUS: deployed
7 REVISION: 1
8 TEST SUITE: None
9 NOTES:
10 1. Get the scheduler address by running these commands:
11   export SCHEDULER_POD_NAME=$(kubectl get pods --namespace dragonfly-syste
  m -l "app=dragonfly,release=dragonfly,component=scheduler" -o jsonpath={.i
  tems[0].metadata.name})
12   export SCHEDULER_CONTAINER_PORT=$(kubectl get pod --namespace dragonfly-
  system $SCHEDULER_POD_NAME -o jsonpath="{.spec.containers[0].ports[0].cont
  ainerPort}")
13   kubectl --namespace dragonfly-system port-forward $SCHEDULER_POD_NAME 80
  02:$SCHEDULER_CONTAINER_PORT
14   echo "Visit http://127.0.0.1:8002 to use your scheduler"
15
16 2. Get the dfdaemon port by running these commands:
17   export DFDAEMON_POD_NAME=$(kubectl get pods --namespace dragonfly-syste
  m -l "app=dragonfly,release=dragonfly,component=dfdaemon" -o jsonpath={.it
  ems[0].metadata.name})
18   export DFDAEMON_CONTAINER_PORT=$(kubectl get pod --namespace dragonfly-s
  ystem $DFDAEMON_POD_NAME -o jsonpath="{.spec.containers[0].ports[0].contai
  nerPort}")
19   You can use $DFDAEMON_CONTAINER_PORT as a proxy port in Node.
20
21 3. Configure runtime to use dragonfly:
22   https://d7y.io/docs/getting-started/quick-start/kubernetes/
23
24
25 4. Get Jaeger query URL by running these commands:
26   export JAEGER_QUERY_PORT=$(kubectl --namespace dragonfly-system get serv
  ices dragonfly-jaeger-query -o jsonpath="{.spec.ports[0].port}")
27   kubectl --namespace dragonfly-system port-forward service/dragonfly-jaeg
  er-query 16686:$JAEGER_QUERY_PORT
28   echo "Visit http://127.0.0.1:16686/search?limit=20&lookback=1h&maxDurati
  on&minDuration&service=dragonfly to query download events"
```

Check that dragonfly is deployed successfully:

```
Shell |
1 $ kubectl get po -n dragonfly-system
2 NAME                                READY   STATUS    RESTARTS   AGE
3 dragonfly-dfdaemon-65rz7            1/1    Running   5 (6m17s ago)  8m4
4   3s
5 dragonfly-dfdaemon-rnvsj            1/1    Running   5 (6m23s ago)  8m4
6   3s
7 dragonfly-jaeger-7d58dfcfc8-qmn8c  1/1    Running   0           8m4
8   3s
9 dragonfly-manager-6f8b4f5c66-qq8sd 1/1    Running   0           8m4
10  3s
11 dragonfly-mysql-0                  1/1    Running   0           8m4
12  3s
13 dragonfly-redis-master-0           1/1    Running   0           8m4
14  3s
15 dragonfly-redis-replicas-0         1/1    Running   0           8m4
16  3s
17 dragonfly-redis-replicas-1         1/1    Running   0           7m3
18  3s
19 dragonfly-redis-replicas-2         1/1    Running   0           5m5
20  0s
21 dragonfly-scheduler-0              1/1    Running   0           8m4
22  3s
23 dragonfly-seed-peer-0              1/1    Running   3 (5m56s ago) 8m4
24  3s
```

Expose Dragonfly Peer Proxy service port

Create the `proxy.yaml` configuration to expose the port on which the Dragonfly Peer Proxy listens. The default port is `65001` and set `targetPort` to `65001`.


```
YAML |
1 kind: Service
2 apiVersion: v1
3 metadata:
4   name: proxy
5 spec:
6   selector:
7     app: dragonfly
8     component: dfdaemon
9     release: dragonfly
10
11   ports:
12   - protocol: TCP
13     port: 65001
14     targetPort: 65001
15
16   type: NodePort
```

Create service:

```
Plain Text |
1 kubectl --namespace dragonfly-system apply -f proxy.yaml
```

Forward request to Dragonfly Peer Proxy:

```
Shell |
1 kubectl --namespace dragonfly-system port-forward service/proxy 65001:65001
```

Install Fluid

Create Fluid cluster based on helm charts

For detailed installation documentation, please refer to [document](#).

Create namespace:

```
Shell |
1 $ kubectl create ns fluid-system
```

Create a dragonfly cluster:

```
Shell |
1 $ helm repo add fluid https://fluid-cloudnative.github.io/charts
2 $ helm repo update
3 $ helm install fluid fluid/fluid
4 NAME: fluid
5 LAST DEPLOYED: Thu Oct 12 21:54:34 2023
6 NAMESPACE: default
7 STATUS: deployed
8 REVISION: 1
9 TEST SUITE: None
```

Check that Fluid is deployed successfully:

```
Shell |
1 $ kubectl get po -n fluid-system
2 NAME                                READY   STATUS    RESTARTS   AGE
3 csi-nodeplugin-fluid-nq65p          2/2    Running   0           7m12s
4 csi-nodeplugin-fluid-nrwbt          2/2    Running   0           7m12s
5 csi-nodeplugin-fluid-q565r          2/2    Running   0           7m12s
6 dataset-controller-5f5f46d969-lpsc7 1/1    Running   0           7m11s
7 fluid-webhook-75f489c7b5-whzjz      1/1    Running   0           7m12s
8 fluidapp-controller-54975849ff-w272h 1/1    Running   0           7m12s
```

Create Dataset

Create the `secret.yaml` :

```
Shell |
1 apiVersion: v1
2 kind: Secret
3 metadata:
4   name: jfs-secret
5 type: Opaque
6 stringData:
7   name: dragonfly
8   metaurl: redis://127.0.0.1:6379/3
9   token: ${JUICEFS_TOKEN}
10  access-key: ${ACCESS_KEY}
11  secret-key: ${SECRET_KEY}
```

Create a secret using the configuration:

```
1 $ kubectl create -f secret.yaml
```

Create `dataset.yaml` :

```
1 $ cat<<EOF >dataset.yaml
2 apiVersion: data.fluid.io/v1alpha1
3 kind: Dataset
4 metadata:
5   name: jfsdemo
6 spec:
7   mounts:
8     - name: dragonfly
9       mountPoint: "juicefs:///"
10      options:
11        bucket: "'https://myjfs.oss-cn-shanghai.aliyuncs.com?proxy=http://
127.0.0.1:65001&backendStorage=oss'"
12        storage: "dragonfly"
13      encryptOptions:
14        - name: metaurl
15          valueFrom:
16            secretKeyRef:
17              name: jfs-secret
18              key: metaurl
19        - name: token
20          valueFrom:
21            secretKeyRef:
22              name: jfs-secret
23              key: token
24        - name: access-key
25          valueFrom:
26            secretKeyRef:
27              name: jfs-secret
28              key: access-key
29        - name: secret-key
30          valueFrom:
31            secretKeyRef:
32              name: jfs-secret
33              key: secret-key
34 EOF
```

Where:

- `mountPoint` : The directory where users store data in the JuiceFS file system, starting

with `juicefs://` . For example, `juicefs:///demo` is a subdirectory `/demo` of the JuiceFS file system.

- `storage` : Set the object storage type, the value is `dragonfly` .
- `bucket` : Set the object storage endpoint.

Dragonfly Peer Proxy address and Backend object storage type are added to the query string in the `--bucket` . The details of parameter is as follows:

Param	Type	Describe	Required
backendStorage	string	Backend object storage type, supports s3, oss and obs.	Y
proxy	string	Draognfly Peer Proxy address.	Y

Create a Dataset:

```
▼ Shell |
1 $ kubectl create -f dataset.yaml
2 dataset.data.fluid.io/jfsdemo created
```

Create JuiceFS Runtime

Create `runtime.yaml` :

```
▼ Shell |
1 $ cat<<EOF >runtime.yaml
2 apiVersion: data.fluid.io/v1alpha1
3 kind: JuiceFSRuntime
4 metadata:
5   name: jfsdemo
6 spec:
7   replicas: 1
8   fuse:
9     image: dragonflyoss/juicefs-fuse
10    imageTag: 0.1.0
11    imagePullPolicy: IfNotPresent
12  juicefsVersion:
13    image: dragonflyoss/juicefs-fuse
14    imageTag: 0.1.0
15    imagePullPolicy: IfNotPresent
16  tieredstore:
17    levels:
18      - mediumtype: MEM
19        path: /dev/shm
20        quota: 40Gi
21        low: "0.1"
22 EOF
```

Create JuiceFS Runtime:

```
▼ Shell |
1 $ kubectl create -f runtime.yaml
2 juicefsruntime.data.fluid.io/jfsdemo created
```

JuiceFS Runtime to start successfully:

```
▼ Shell |
1 $ kubectl get po |grep jfs
2 jfsdemo-worker-0 1/1 Running
   0 4m2s
```

Check the dataset status, it has been bound to JuiceFS Runtime:

```

Shell |
1 $ kubectl get dataset jfsdemo
2 NAME          UFS TOTAL SIZE    CACHED    CACHE CAPACITY    CACHED PERCENTAGE    PH
   ASE    AGE
3 jfsdemo      0.00B            0.00B    40.00GiB        0.0%                Bo
   und    21h

```

Fluid has created PV and PVC with the same name as the dataset.

```

Shell |
1 $ kubectl get pv | grep jfs
2 default-jfsdemo      100Pi      ROX      Retain      Bound      d
   default/jfsdemo
   16h      fluid
3 $ kubectl get pvc
4 NAME          STATUS    VOLUME          CAPACITY    ACCESS MODES    STORAGECLAS
   S    AGE
5 jfsdemo      Bound    default-jfsdemo  100Pi      ROX              fluid
   21h

```

Verify

Create an application to use the dataset:

```
Shell |
1 $ cat<<EOF >app.yaml
2   apiVersion: v1
3   kind: Pod
4   metadata:
5     name: demo-app
6   spec:
7     containers:
8     - name: demo
9       image: nginx
10      volumeMounts:
11      - mountPath: /demo
12        name: demo
13     volumes:
14     - name: demo
15       persistentVolumeClaim:
16         claimName: jfsdemo
17 EOF
```

Create application:

```
Shell |
1 $ kubectl create -f app.yaml
```

Check that the Pod has been created:

```
Shell |
1 $ kubectl get pod
2 NAME                                READY   STATUS    RESTARTS   AGE
3 demo-app                             1/1     Running   0           40s
4 jfsdemo-fuse-vfqgt                   1/1     Running   0           40s
5 jfsdemo-worker-0                      1/1     Running   0           102s
```

The Pod has been created successfully and JuiceFS's FUSE component has also started successfully.

Reference

1. [JuiceFS Install](#)
2. [JuiceFS Performance Testing](#)
3. [JuiceFS Command Reference](#)

4. [Dragonfly Quick Start](#)
5. [Dragonfly Helm chart configuration file](#)
6. [How to Use JuiceFS in Fluid](#)